

Table of Contents

1. Introduction.....	2
1.1 First Recorded Use.....	2
1.2 Use Within Business.....	3
1.3 Description of dataset.....	3
2. Packages.....	4
2.1 Required R – Packages.....	4
2.1 Required Python 3 – Packages.....	5
2.3 Which is better for data analysis: R or Python.....	6
3. Exploratory Data Analysis (EDA).....	8
3.1 Data Structure.....	7
4. Exploratory Data Analysis (EDA) in R.....	11
4.1 Gender.....	13
4.2 Top Sellers.....	18
4.3 Age.....	22
4.4 Purchase.....	27
4.5 Marital Status.....	30
4.6 Top Shoppers.....	37
4.7 Occupation.....	39
4.8 Apriori (Association Rule Learning).....	41
5. Conclussions.....	55
6. References.....	56
7. Packages.....	56
8. Images.....	57
9. Tables.....	57
10. Figures.....	58

1. Introduction

The origins of "Black Friday" stem not from a day filled with shopping, discounts, and a turn of the holiday season, but rather with a financial crisis. The first recorded use of the term "Black Friday" was recorded on September 24th, 1869 when two Wall Street businessmen, Jay Gould and Jim Fisk, decided to artificially inflate the price of gold and attempted to sell it for profit. The result of their nefarious actions, on that specific Friday in September 1869, the price of Gold dropped and the United States plunged into a state of financial devastation.¹



Picture 1. <https://justcreative.com/wp-content/uploads/2018/10/black-friday-deals.jpg>

1.1 First Recorded Use

Various stories exist regarding the first recorded use of the term as it relates to holiday shopping, but its connotation continued to keep a negative stigma associated with it until the late 20th century. "Black Friday" and its relation to consumerism first derived from 1950s Philadelphia. Philadelphia suburbanites descended on the city after the Thanksgiving holiday, to watch the traditional Army college football game and take advantage of sales and promotions, Philadelphia Police Officers who were assigned to work that weekend coined the term due to their long grueling shifts and the mass amounts of people/shoppers. Philadelphia businesses also started to use the term to describe the long lines and shopping mayhem at their stores.

¹ [https://en.wikipedia.org/wiki/Black_Friday_\(shopping\)](https://en.wikipedia.org/wiki/Black_Friday_(shopping))

1.2. Use Within Business

One of the possible explanation for the term as it relates to consumers and retailers is that "Black Friday" represents the first day of the year in which businesses were turning profits and accounting was done on a hand-written ledger. As described Oxford Dictionary, "The use of colors in accounting refers back to the bookkeeping practice of recording the credit side of an account in a ledger in *black* ink and the debit side in *red* ink." (Oxford Dictionaries) Hence the name, "Black Friday" being associated with businesses debits overtaking their credits. Although this idea might make sense, the claim hasn't been completely verified.²

1.3 Description of dataset

The dataset here is a sample of the transactions made in a retail store. The store wants to know better the customer purchase behaviour against different products. Specifically, here the problem is a regression problem where we are trying to predict the dependent variable (the amount of purchase) with the help of the information contained in the other variables.

Classification problem can also be settled in this dataset since several variables are categorical, and some other approaches could be "Predicting the age of the consumer" or even "Predict the category of goods bought". This dataset is also particularly convenient for clustering and maybe find different clusters of consumers within it³

1. Main aims of the project are:

- analysis all of dataset in relation to variables
- showing the possibility of packages in R and the Python

2. My research hypotheses are:

1. The male customers have a higher average spending then the female.
2. Which gender and in which age have achavied more purchase.
3. The customers are young people.

² Oxford Dictionaries- Black Friday

³ <https://www.kaggle.com/mehdidag/black-friday>

2. Required R – Packages

There are several R packages that useful for analyzing this dataset.

- dplyr - tool frome processing dataset. 2.1
- ggplot2 - creating graphics. 2.2
- plotly - to help make pie chart. 2.3

2.1.The dplyr package makes these steps fast and easy:

- By constraining your options, it helps you think about your data manipulation challenges.
- It provides simple “verbs”, functions that correspond to the most common data manipulation tasks, to help you translate your thoughts into code.
- It uses efficient backends, so you spend less time waiting for the computer.⁴

2.2 The ggplot2 is a data visualization package for the statistical programming language R. Created by Hadley Wickham in 2005, ggplot2 is an implementation of Leland Wilkinson's *Grammar of Graphics* — a general scheme for data visualization which breaks up graphs into semantic components such as scales and layers. ggplot2 can serve as a replacement for the base graphics in R and contains a number of defaults for web and print display of common scales. Since 2005, ggplot2 has grown in use to become one of the most popular R packages.⁵

2.3 Plotly's R graphing library makes interactive, publication-quality graphs online. Examples of how to make line plots, scatter plots, area charts, bar charts, error bars, box plots, histograms, heatmaps, subplots, multiple-axes, and 3D (WebGL based) charts.⁶

⁴ <https://cran.r-project.org/web/packages/dplyr/vignettes/dplyr.html>

⁵ <https://en.wikipedia.org/wiki/Ggplot2>

⁶ <https://plot.ly/r/>

2.4 Required Python 3 – Packages

numpy as np - linear algebra

pandas as pd - data processing, CSV file I/O (e.g. pd.read_csv)

matplotlib.pyplot as plt

There are several Python packages that useful for analyzing dataset's:

2.4.1. NumPy is the fundamental package for scientific computing with Python. It contains among other things:⁷

- a powerful N-dimensional array object
- sophisticated (broadcasting) functions
- tools for integrating C/C++ and Fortran code
- useful linear algebra, Fourier transform, and random number capabilities

2.4.2. Pandas is a Python package providing fast, flexible, and expressive data structures designed to make working with “relational” or “labeled” data both easy and intuitive. It aims to be the fundamental high-level building block for doing practical, real world data analysis in Python. Additionally, it has the broader goal of becoming the most powerful and flexible open source data analysis and manipulation tool available in any language.⁸

Pandas is well suited for many different kinds of data:

- Tabular data with heterogeneously - typed columns, as in an SQL table or Excel spreadsheet
- Ordered and unordered (not necessarily fixed-frequency) time series data.
- Arbitrary matrix data (homogeneously typed or heterogeneous) with row and column labels
- Any other form of observational / statistical data sets. The data actually need not be labeled at all to be placed into a pandas data structure.

2.4.3. Matplotlib is a Python 2D plotting library which produces publication quality figures in a variety of hardcopy formats and interactive environments across platforms. Matplotlib can be used in Python scripts, the Python and IPython shells, the Jupyter notebook, web application servers, and four graphical user interface toolkits.⁹

⁷ <https://www.numpy.org/>

⁸ <https://pandas.pydata.org/>

⁹ <https://matplotlib.org/>

2.3 Which is better for data analysis: R or Python

R is considered to be the best programming language for any statistician as it possesses an extensive catalog of statistical and graphical methods. Python on the other hand can do pretty much the same work as R but it is preferred by the data scientists or data analysts because of its simplicity and high performance. R is a powerful scripting language and highly flexible with a vibrant community and resource bank whereas Python is a widely used, object oriented language which is easy to learn and debug.

R has a much bigger library of statistical packages

If we doing specialized statistical work, R packages cover more techniques. We can find R packages for a wide variety of statistical tasks using the CRAN task view. R packages cover everything from Psychometrics to Genetics to Finance. Although Python, through SciPy and packages like statsmodels, covers the most common techniques, R is far ahead.

Python is better for building analytics tools

R and Python are equally good if you want to find outliers in a dataset, but if we want to create a web service to enable other people to upload datasets and find outliers, Python is better. Python is a general purpose programming language, which means that people have built modules to create websites, interact with a variety of databases, and manage users.

R builds in data analysis functionality by default, whereas Python relies on packages

Python is a general purpose language, most data analysis functionality is available through packages like NumPy and pandas. However, R was built with statistics and data analysis in mind, so many tools that have been added to Python through packages are built into base R.

Python is better for deep learning

Through packages like Lasagne, caffe, keras, and tensorflow, creating deep neural networks is straightforward in Python. Although some of these, like tensorflow, are being ported to R, support is still far better in Python.

Python relies on a few main packages, whereas R has hundreds

In Python, sklearn is the “primary” machine learning package, and pandas is the “primary” data analysis package. This makes it easy to know how to accomplish a task, but also means that a lot of specialized techniques aren’t possible.

R is better for data visualization

Packages like ggplot2 make plotting easier and more customizable in R than in Python. Python is catching up, particularly in the area of interactive plots with packages like Bokeh.

3.1 Exploratory Data Analysis (EDA)

To load the dataset that we will be using for this Exploratory Data Analysis (EDA).

```
#importing the dataset
data = pd.read_csv('../input/BlackFriday.csv')
This dataset has 12 variables
```

##	User_ID	Product_ID	Gender	Age	Occupation	City_Category
## 1	1000001	P00069042	F	0-17	10	A
## 2	1000001	P00248942	F	0-17	10	A
## 3	1000001	P00087842	F	0-17	10	A
## 4	1000001	P00085442	F	0-17	10	A
## 5	1000002	P00285442	M	55+	16	C
## 6	1000003	P00193542	M	26-35	15	A
##	Stay_In_Current_City_Years			Marital_Status	Product_Category_1	
## 1	2			0	3	
## 2	2			0	1	
## 3	2			0	12	
## 4	2			0	12	
## 5	4+			0	8	
## 6	3			0	1	
##	Product_Category_2		Product_Category_3		Purchase	
## 1	NA		NA		8370	
## 2	6		14		15200	
## 3	NA		NA		1422	
## 4	14		NA		1057	
## 5	NA		NA		7969	
## 6	2		NA		15227	

Table 1. Code from R program with data analyses.
Source: Own elaboration.

To read and analyze data by using Pandas library:

We will also import all libraries which will be used.

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt
import seaborn as sns # visualization tool
import os
print(os.listdir("../input"))

# read BlackFriday.csv file data from input directory and create dataframe named data
data = pd.read_csv("../input/BlackFriday.csv")
```

After reading csv file , we will check data by using „info()” method of dataframe , we can see column names and they types .

When we execute command we will see there are 537577 entries in file. For Product_Category_1 column , 164278 of 537577 are non-null and that means rest are null. For Product_Category_2 column , 370591 of 537577 are non-null and that means rest are null and in other columns are all full and there is no empty data.

```
# check information about columns data
data.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 537577 entries, 0 to 537576
Data columns (total 12 columns):
User_ID                537577 non-null int64
Product_ID             537577 non-null object
Gender                 537577 non-null object
Age                    537577 non-null object
Occupation              537577 non-null int64
City_Category          537577 non-null object
Stay_In_Current_City_Years  537577 non-null object
Marital_Status         537577 non-null int64
Product_Category_1     537577 non-null int64
Product_Category_2     370591 non-null float64
Product_Category_3     164278 non-null float64
Purchase               537577 non-null int64
dtypes: float64(2), int64(5), object(5)
memory usage: 49.2+ MB
```

Table 2. Information about columns data.
Source: Own elaboration.

When we look on first 10 data of dataframe to have knowledge data itself.

```
#data.head(10)
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3	Purchase
0	1000001	P00069042	F	0-17	10	A	2	0	3	NaN	NaN	8370
1	1000001	P00248942	F	0-17	10	A	2	0	1	6.0	14.0	15200
2	1000001	P00087842	F	0-17	10	A	2	0	12	NaN	NaN	1422
3	1000001	P00085442	F	0-17	10	A	2	0	12	14.0	NaN	1057
4	1000002	P00285442	M	55+	16	C	4+	0	8	NaN	NaN	7969
5	1000003	P00193542	M	26-35	15	A	3	0	1	2.0	NaN	15227
6	1000004	P00184942	M	46-50	7	B	2	1	1	8.0	17.0	19215
7	1000004	P00346142	M	46-50	7	B	2	1	1	15.0	NaN	15854
8	1000004	P0097242	M	46-50	7	B	2	1	1	16.0	NaN	15686
9	1000005	P00274942	M	26-35	20	A	1	1	8	NaN	NaN	7871

Table 3. First 10 data of dataframe.
Source: Own elaboration.

We can see there are **NaN** m_values for Product_Category_2 and Product_Category_3.

We want to fill them with zero.

```
data.Product_Category_2.fillna(0, inplace=True)
data.Product_Category_3.fillna(0, inplace=True)
```

By using describe method of **dataFrame** , we can learn some statistical information such as mean(average), max, min etc about data.

```
# display columns
data.columns
Index(['User_ID', 'Product_ID', 'Gender', 'Age', 'Occupation',
       'City_Category',
       'Stay_In_Current_City_Years', 'Marital_Status',
       'Product_Category_1',
       'Product_Category_2', 'Product_Category_3', 'Purchase'],
      dtype='object')
```

Now we can check statistical information about numeric data columns

```
data.describe()
```

	User_ID	Occupation	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3	Purchase
count	5.375770e+05	537577.00000	537577.00000	537577.000000	537577.000000	537577.000000	537577.000000
Mean	1.002992e+06	8.08271	0.408797	5.295546	6.784907	3.871773	9333.859853
std	1.714393e+03	6.52412	0.491612	3.750701	6.211618	6.265963	4981.022133
min	1.000001e+06	0.00000	0.000000	1.000000	0.000000	0.000000	185.000000
25%	1.001495e+06	2.00000	0.000000	1.000000	0.000000	0.000000	5866.000000
50%	1.003031e+06	7.00000	0.000000	5.000000	5.000000	0.000000	8062.000000
75%	1.004417e+06	14.00000	1.000000	8.000000	14.000000	8.000000	12073.000000
max	1.006040e+06	20.00000	1.000000	18.000000	18.000000	18.000000	23961.000000

Table 4 . First 10 data of dataframe to have knowledge data itself.
Source: Own elaboration.

```
data.corr()
```

	User_ID	Occupation	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3	Purchase
User_ID	1.000000	-0.023024	0.018732	0.003687	0.003663	0.003938	0.005389
Occupation	-0.023024	1.000000	0.024691	-0.008114	0.006792	0.011941	0.021104
Marital_Status	0.018732	0.024691	1.000000	0.020546	0.001146	-0.004363	0.000129
Product_Category_1	0.003687	-0.008114	0.020546	1.000000	-0.040730	-0.389048	-0.0314125
Product_Category_2	0.003663	0.006792	0.001146	-0.040730	1.000000	0.090284	0.038395
Product_Category_3	0.003938	0.011941	-0.004363	-0.389048	0.090284	1.000000	0.284120
Purchase	0.005389	0.021104	0.000129	-0.314125	0.038395	0.284120	1.000000

Table 5. First 7 data of dataframe to have knowledge data itself.

Source: Own elaboration

4. Exploratory Data Analysis (EDA) in R

The tidyverse package is what we will use for visualizing and exploring our dataset.

It is known that for easy to read syntax and massive amounts of useful functions. The scales package will be used mainly to customize plot axis. Lastly the arules package will be utilized in the final part of this kernel, Association Rule Learning and Apriori. Info regarding all packages used during this EDA is provided in the Works Cited section in this kernel.

Lets start with overview of the entire dataset.

User_ID	Product_ID	Gender	Age
Min. :1000001	P00265242:	1858 F:132197	0-17 : 14707
1st Qu.:1001495	P00110742:	1591 M:405380	18-25: 97634
Median :1003031	P00025442:	1586	26-35:214690
Mean :1002992	P00112142:	1539	36-45:107499
3rd Qu.:1004417	P00057642:	1430	46-50: 44526
Max. :1006040	P00184942:	1424	51-55: 37618
(Other) :528149		55+ : 20903	
Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status
Min. : 0.000	A:144638	0 : 72725	Min. :0.0000
1st Qu.: 2.000	B:226493	1 : 189192	1st Qu.:0.0000
Median : 7.000	C:166446	2 : 99459	Median :0.0000
Mean : 8.083		3 : 93312	Mean :0.4088
3rd Qu.:14.000		4+: 82889	3rd Qu.:1.0000
Max. :20.000			Max. :1.0000

Product_Category_1	Product_Category_2	Product_Category_3	Purchase
Min. : 1.000	Min. : 2.00	Min. : 3.0	Min. : 185
1st Qu.: 1.000	1st Qu.: 5.00	1st Qu.: 9.0	1st Qu.: 5866
Median : 5.000	Median : 9.00	Median :14.0	Median : 8062
Mean : 5.296	Mean : 9.84	Mean :12.7	Mean : 9334
3rd Qu.: 8.000	3rd Qu.:15.00	3rd Qu.:16.0	3rd Qu.:12073
Max. :18.000	Max. :18.00	Max. :18.0	Max. :23961
NA's :166986	NA's :373299		

Table 6 . Overview of the entire dataset.
Source: Own elaboration.

User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3	Purchase
1000001	P00069042	F	0-17	10	A	2	0	3	NA	NA	8370
1000001	P00248942	F	0-17	10	A	2	0	1	6	14	15200
1000001	P00087842	F	0-17	10	A	2	0	12	NA	NA	1422
1000001	P00085442	F	0-17	10	A	2	0	12	14	NA	1057
1000002	P00285442	M	55+	16	C	4+	0	8	NA	NA	7969
1000003	P00193542	M	26-35	15	A	3	0	1	2	NA	15227

Table 7 . Overview of the entire dataset.
Source: Own elaboration.

We have 12 different columns, each representing a corresponding variable below.

- User_ID: Unique identifier of shopper.
- Product_ID: Unique identifier of product. (No key given)
- Gender: Sex of shopper.
- Age: Age of shopper split into bins.
- Occupation: Occupation of shopper. (No key given)
- City_Category: Residence location of shopper. (No key given)

- Stay_In_Current_City_Years: Number of years stay in current city.
- Marital_Status: Marital status of shopper.
- Product_Category_1: Product category of purchase.
- Product_Category_2: Product may belong to other category.
- Product_Category_3: Product may belong to other category.
- Purchase: Purchase amount in dollars.

If we look at the first few rows of our dataset, we can see that each row represents a different transaction or item purchased by a specific customer. When we group all transactions by a specific User_ID to get a sum of all purchases made by a single customer.

One critique we can make regarding this dataset is that there isn't a key given regarding the different Product_IDs and the item they represent. (Ie. We can't attribute P00265242 to an item easily recognizable) In reality, we would want to have another dataset which provides the name of an Item and its Product_ID and then join it to our existing dataset. This won't necessarily affect our EDA, but would be more useful during our implementation of the Apriori algorithm and could make some parts of the EDA clearer to interpret.

4.1 Gender

To begin our exploration lets examine the gender of shoppers at this store.

Since each row represents an individual transaction, we must first group the data by User_ID to remove duplicates.

```
dataset_gender = dataset %>%
  select(User_ID, Gender) %>%
  group_by(User_ID) %>%
  distinct()

head(dataset_gender)

summary(dataset_gender$Gender)
```

User_ID	Gender
1000001	F
1000002	M
1000003	M
1000004	M
1000005	M
1000006	F

Table 8. Duplctes.
Source: Own elaboration.

F	1666
M	4225

Table 9. F and M in data Duplictcs.
Source: Own elaboration.

We have the dataframe necessary to see each User_IDs corresponding gender and their total counts for reference, lets plot the distribution of gender across our dataset.

```
options(scipen=10000) # To remove scientific numbering

genderDist = ggplot(data = dataset_gender) +
  geom_bar(mapping = aes(x = Gender, y = ..count.., fill = Gender)) +
  labs(title = 'Gender of Customers') +
  scale_fill_brewer(palette = 'PuBuGn')
print(genderDist)
```




Figure 2. Gender plot.
Source: Own elaboration.

We can see there are quite a few more males than females shopping at our store on Black Friday. This gender split metric is helpful to retailers because some might want to modify their store layout product selection, and other variables differently depending on the gender proportion of their shoppers. (Figure 2)

When we study published in the Clothing and Textiles Research Journal writes,

- "Involvement, variety seeking, and physical environment of stores were selected as antecedents of shopping experience satisfaction....The structural model for female subjects confirmed the existence of the mediating role of hedonic shopping value in shopping satisfaction, whereas the model for male respondents did not." Chang, E., Burns, L. D., & Francis, S. K. (2004) (Abstract)¹⁰

Although this does not give direct insight into recommended actions for retail stores, it does display a difference in the value derived from shopping and its relationship to gender, which should be taken into account by retailers.

To investigate further, lets compute the average spending amount as it relates to Gender.

For easy interpretation and traceback we will create separate tables and then join them together.

```
total_purchase_user = dataset %>%
  select(User_ID, Gender, Purchase) %>%
  group_by(User_ID) %>%
  arrange(User_ID) %>%
  summarise(Total_Purchase = sum(Purchase))

user_gender = dataset %>%
  select(User_ID, Gender) %>%
  group_by(User_ID) %>%
  arrange(User_ID) %>%
  distinct()

head(user_gender)
head(total_purchase_user)
```

User_ID	Gender
1000001	F
1000002	M
1000003	M
1000004	M
1000005	M
1000006	F

Table 11. Total Purchase
Source: Own elaboration

¹⁰ Chang, E., Burns, L. D., & Francis, S. K. (2004) (Abstract)

User_ID	Total_Purchase
1000001	333481
1000002	810353
1000003	341635
1000004	205987
1000005	821001
1000006	379450

Table 11a. Total Purchase
Source: Own elaboration.

```
user_purchase_gender = full_join(total_purchase_user, user_gender, by
= "User_ID")
head(user_purchase_gender)
```

User_ID	Total_Purchase	Gender
1000001	333481	F
1000002	810353	M
1000003	341635	M
1000004	205987	M
1000005	821001	M
1000006	379450	F

Table 12. Join Table 2 and Table 3 together.
Source: Own elaboration.

```
average_spending_gender = user_purchase_gender %>%
  group_by(Gender) %>%
  summarize(Purchase = sum(as.numeric(Total_Purchase)),
            Count = n(),
            Average = Purchase/Count)
head(average_spending_gender)
```

Gender	Purchase	Count	Average
F	1164624021	1666	699054.0
M	3853044357	4225	911963.2

Table 13. Join Purchase and Average Together
Source: Own elaboration.

```
genderAverage = ggplot(data = average_spending_gender) +
  geom_bar(mapping = aes(x = Gender, y = Average, fill = Gender),
stat = 'identity') +
```



```
      labs(title = 'Average Spending by Gender') +  
      scale_fill_brewer(palette = 'PuBuGn')  
print(genderAverage)
```

Looks like our top 5 best sellers are (by product ID)

- P00265242 = 1858
- P00110742 = 1591
- P00025442 = 1586
- P00112142 = 1539
- P00057642 = 1430

When we have Identified our top 5 best selling products, lets examine the best selling product, P00265242.

```
best_seller = dataset[dataset$Product_ID == 'P00265242', ]
```

```
head(best_seller)
```

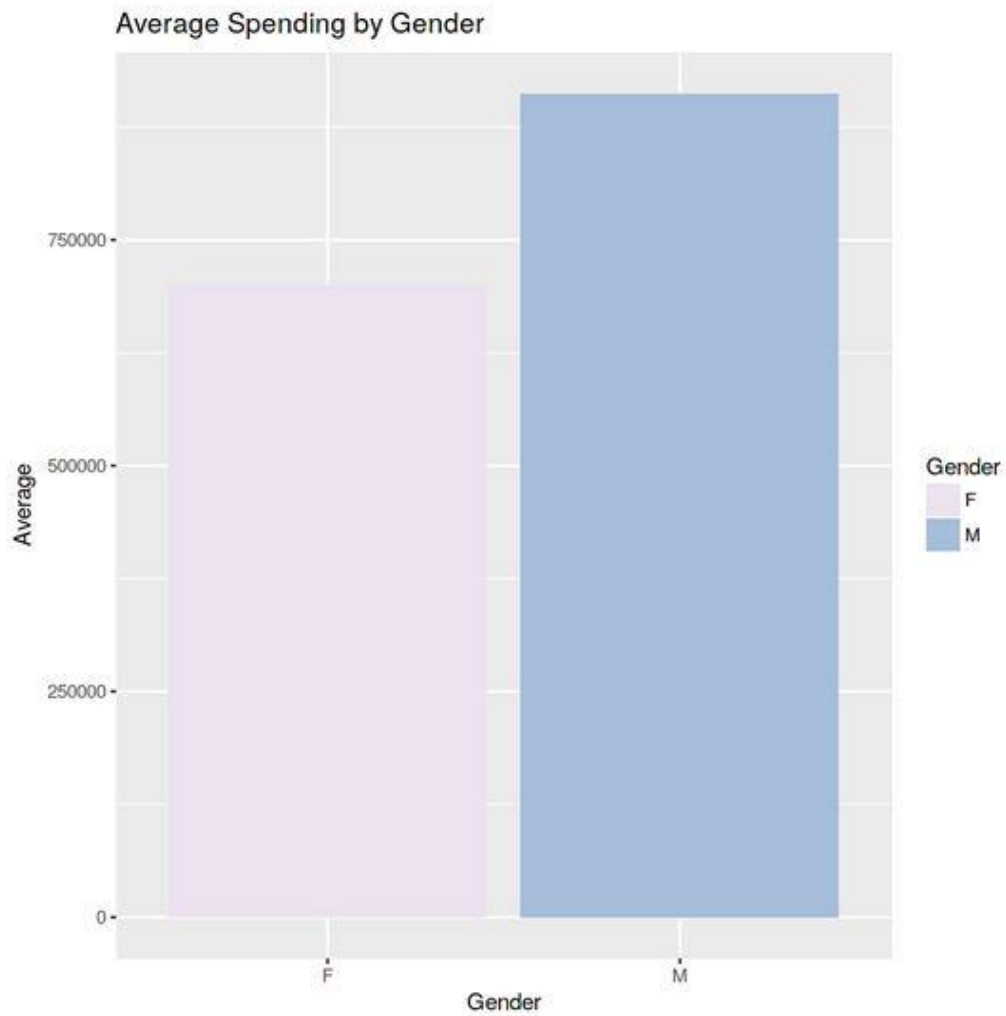


Figure 3. Average spending by Gender plot.
Source: Own elaboration.

We present's an interesting observation. Even though female shoppers make less purchases than males at this specific store, they seem to be purchasing almost as much on average as the male shoppers. This being said, scale needs to be taken into account because females on average are still spending about 250,000 less than males. [Figure 3]

4.2 Top Sellers

Now lets switch gears and examine our top selling products. In this situation, we won't group by product ID since we want to see duplicates, just in case people are buying 2 or more quantities of the same product.

```
top_sellers = dataset %>%  
  count(Product_ID, sort = TRUE)  
  
top_5 = head(top_sellers, 5)  
  
top_5
```

Looks like our top 5 best sellers are (by product ID)

- P00265242 = 1858
- P00110742 = 1591
- P00025442 = 1586
- P00112142 = 1539
- P00057642 = 1430

Now that we have Identified our top 5 best selling products, lets examine the best selling product, P00265242.

	Use r_I D	Prod uct_I D	Ge nd er	A ge	Occu pation	City_ Catego ry	Stay_In_Curr ent_City_Yea rs	Marita l_Statu s	Product_ Category _1	Product_ Category _2	Product_ Category _3	Pur cha se
4 0 0	100 006 6	P002 6524 2	M	2 6 - 3 5	18	C	2	0	5	8	NA	865 2
1 1 9 2	100 019 6	P002 6524 2	F	3 6 - 4 5	9	C	4+	0	5	8	NA	876 7
1 3 7 3	100 022 2	P002 6524 2	M	2 6 - 3 5	1	A	1	0	5	8	NA	694 4
1 8 4 6	100 030 1	P002 6524 2	M	1 8 - 2 5	4	B	4+	0	5	8	NA	862 8
2 2 1 0	100 034 5	P002 6524 2	M	2 6 - 3 5	12	A	2	1	5	8	NA	859 3
2 4 0 5	100 038 3	P002 6524 2	F	2 6 - 3 5	7	A	4+	1	5	8	NA	699 8

Table 13. 5 best selling products.
Source: Own elaboration.

We can see that this product fits into Product_Category_1 = 5 and Product_Category_2 = 8.

As mentioned in the introduction, it would be useful to have a key to reference the item name in order to determine what it is. [Table 13]

Another interesting finding is that even though people are purchasing the same product they are paying different prices.

This could be due to various Black Friday promotions, discounts, or coupon codes. Otherwise, investigation would need to be done regarding the reason for different purchase prices of the same product between customers.

Lets continue to analyze our best seller to see if any relationship to Gender exists.

```
genderDist_bs = ggplot(data = best_seller) +
  geom_bar(mapping = aes(x = Gender, y = ..count.., fill = Gender))
+
  labs(title = 'Gender of Customers (Best Seller)') +
  scale_fill_brewer(palette = 'PuBuGn')
print(genderDist_bs)
```

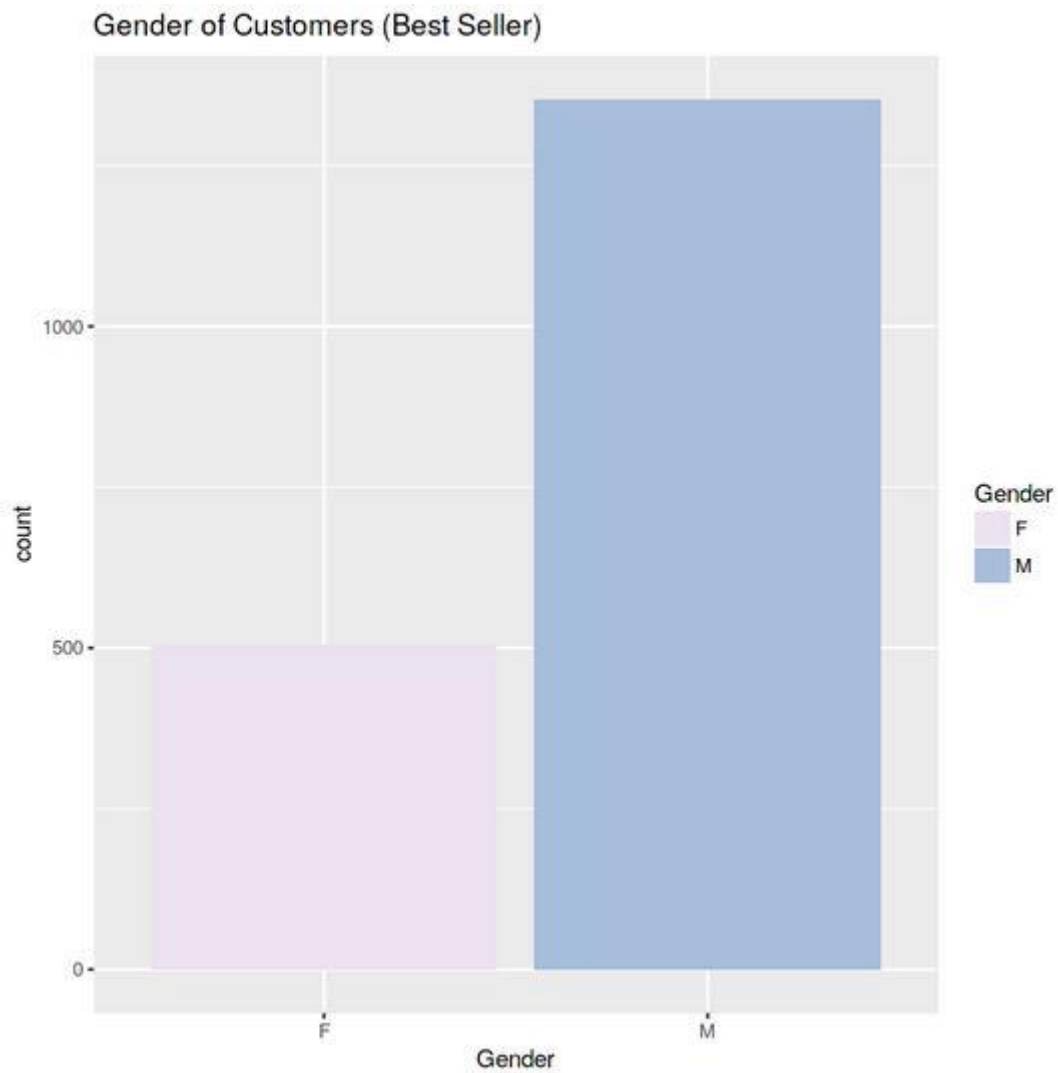


Figure 4. Count and Gender plot
Source: Own elaboration.

We see a similar distribution between genders to our overall dataset gender split - lets confirm.

```
genderDist_bs_prop = ggplot(data = best_seller) +
  geom_bar(fill = 'lightblue', mapping = aes(x = Gender, y =
    ..prop..., group = 1, fill = Gender)) +
  labs(title = 'Gender of Customers (Best Seller -
    Proportion)') +
  theme(plot.title = element_text(size=9.5))

genderDist_prop = ggplot(data = dataset_gender) +
  geom_bar(fill = "lightblue4", mapping = aes(x = Gender, y =
    ..prop..., group = 1)) +
  labs(title = 'Gender of Customers (Total Proportion)') +
  theme(plot.title = element_text(size=9.5))

grid.arrange(genderDist_prop, genderDist_bs_prop, ncol=2)
```




Figure 5. Proportion of customers




Figure 6. Proportion of customers

Source: Own elaboration.

We can see that between the overall observation set, both purchasers of the best seller and purchasers of all products are roughly ~25% female and ~75% male. A slight difference does exist but it seems like we can generally conclude that our best seller does not cater to a specific gender. [Figure 5 and Figure 6]

4.3 Age

Lets begin examining Age by creating a table of each individual age group and their respective counts.
[Table 13]

```
customers_age = dataset %>%  
  select(User_ID, Age) %>%  
  distinct() %>%  
  count(Age)  
  
customers_age
```

Age	n
0-17	218
18-25	1069
26-35	2053
36-45	1167
46-50	531
51-55	481
55+	372

Table 13. Age
Source: Own elaboration.

We can see a dataset that shows the count of each Age category of customers at our store.

```
customers_age_vis = ggplot(data = customers_age) +  
  geom_bar(color = 'black', stat = 'identity',  
mapping = aes(x = Age, y = n, fill = Age)) +  
  labs(title = 'Age of Customers') +  
  theme(axis.text.x = element_text(size = 10)) +  
  scale_fill_brewer(palette = 'Blues') +  
  theme(legend.position="none")  
  
print(customers_age_vis)
```




Figure 7. Age of customers
Source: Own elaboration.

We can also plot a similar chart depicting the distribution of age within our "best seller" category. This will show us if there is a specific age category that purchased the best selling product more than other shoppers. [Figure 7]

```
ageDist_bs = ggplot(data = best_seller) +
  geom_bar(color = 'black', mapping = aes(x = Age, y
= ..count.., fill = Age)) +
  labs(title = 'Age of Customers (Best Seller)') +
  theme(axis.text.x = element_text(size = 10)) +
  scale_fill_brewer(palette = 'GnBu') +
  theme(legend.position="none")
print(ageDist_bs)
```

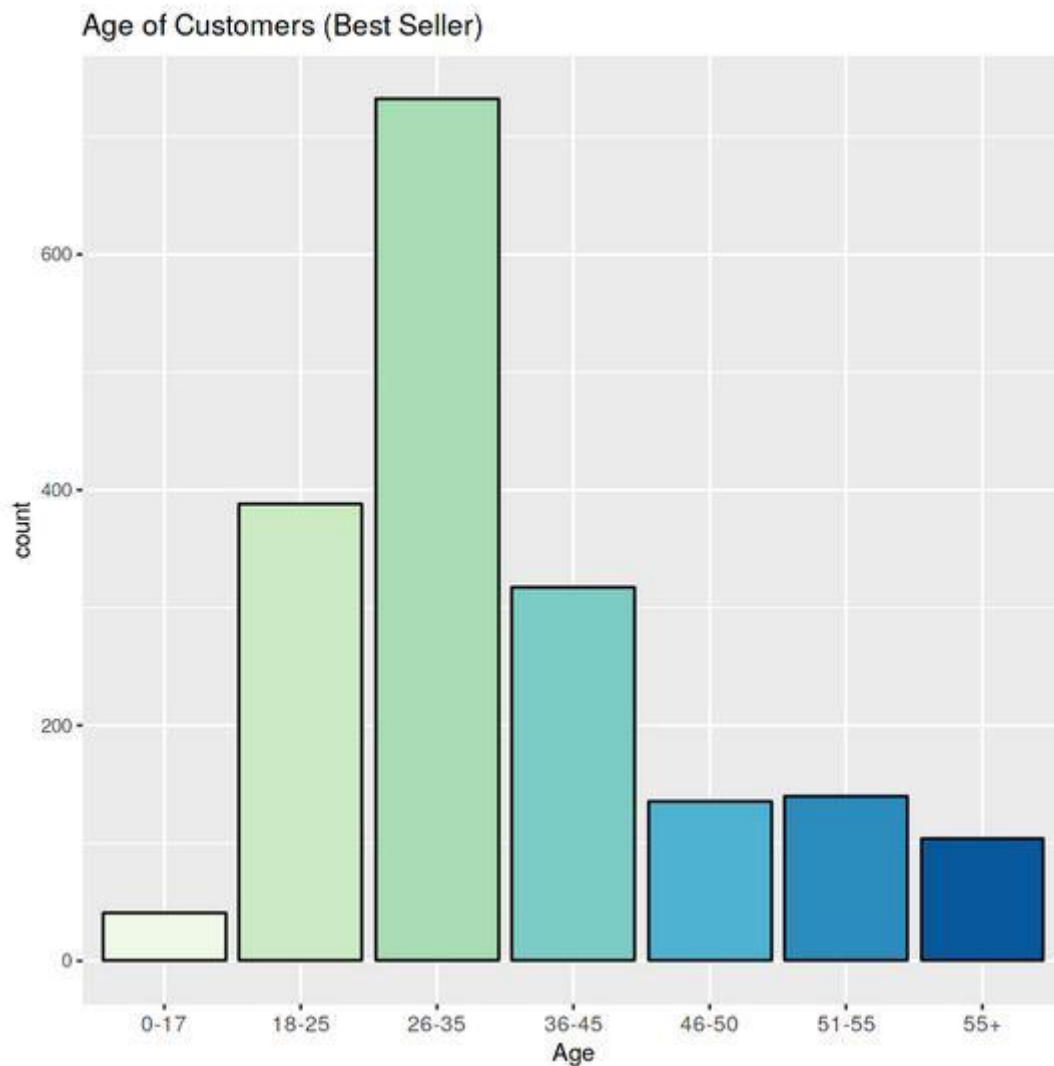



Figure 8. Age of customers best seller.
Source: Own elaboration.

It seems as though younger people (18-25 & 26-35) account for the highest number of purchases of the best selling product. Lets compare this observation to the overall dataset. [Figure 8]

```
grid.arrange(customers_age_vis, ageDist_bs, ncol=2)
```




Figure 9. Age of customers best seler




Figure 10. Age of customers best seller

Source: Own elaboration.

We can see that there is some deviation with the proportion of customers grouped by age when comparing the best selling product to the overall dataset. It looks like older customers > Age 45 are buying the top seller slightly less than other products included in the overall dataset. [Figure 9, Figure 10]

4.4 Purchase

Now lets do some investigation regarding store customers and their purchases. We will start by computing the total purchase amount by user ID

```
customers_total_purchase_amount = dataset %>%  
  group_by(User_ID) %>%  
  summarise(Purchase_Amount = sum(Purchase))  
  
head(customers_total_purchase_amount)
```

User_ID	Purchase_Amount
1000001	333481
1000002	810353
1000003	341635
1000004	205987
1000005	821001
1000006	379450

Table 14. Purchase Amount

Source: Own elaboration.

Now that we have grouped our purchases and grouped by User ID, we will sort and find our top spenders.
[Table 14]

```
customers_total_purchase_amount = arrange(customers_total_purchase_amount,  
  desc((Purchase_Amount)))  
  
head(customers_total_purchase_amount)
```

User_ID	Purchase_Amount
1004277	10536783
1001680	8699232
1002909	7577505
1001941	6817493
1000424	6573609
1004448	6565878

Table 15. Purchase Amount

Source: Own elaboration.

Looks like User ID 1004277 is our top spender. Lets use summary() to see other facets of our total customer spending data. [Table 15]

summary(customers_total_purchase_amount)			
User_ID		Purchase_Amount	
Min.	:1000001	Min.	: 44108
1st Qu.	:1001518	1st Qu.	: 234914
Median	:1003026	Median	: 512612
Mean	:1003025	Mean	: 851752
3rd Qu.	:1004532	3rd Qu.	: 1099005
Max.	:1006040	Max.	:10536783

Table 16. Customers total purchase amount

Source: Own elaboration.

We can see an average total purchase amount of 851752, max total purchase amount of 10536783, min total purchase amount of 44108 and a median purchase amount of 512612.[Table 16]

Lets plot a chart showing the distribution of purchase amounts to see if purchases are normally distributed or contain some skewness. A density plot will show us where the highest number of similar purchase amounts rests in accordance to the entire customer base. It is important to note that Density charts graph the expected probability of values, given data as input, and then plot a line surrounding those values (estimation).

```
ggplot(customers_total_purchase_amount, aes(Purchase_Amount)) +
  geom_density(adjust = 1) +
```

```

geom_vline(aes(xintercept=median(Purchase_Amount)),
  color="blue", linetype="dashed", size=1) +
geom_vline(aes(xintercept=mean(Purchase_Amount)),
  color="red", linetype="dashed", size=1) +
geom_text(aes(x=mean(Purchase_Amount), label=round(mean(Purchase_Amount)), y=1.2e-
06), color = 'red', angle=360,
  size=4, vjust=3, hjust=-.1) +
geom_text(aes(x=median(Purchase_Amount), label=round(median(Purchase_Amount)),
y=1.2e-06), color = 'blue', angle=360,
  size=4, vjust=0, hjust=-.1) +
  scale_x_continuous(name="Purchase Amount", limits=c(0, 7500000), breaks =
seq(0,7500000, by = 1000000), expand = c(0,0)) +
  scale_y_continuous(name="Density", limits=c(0, .00000125), labels = scientific,
expand = c(0,0))
Warning message:
"Removed 3 rows containing non-finite values (stat_density)."
```




Figure 11. Purchase Amount to Density

Source: Own elaboration.

We see a very right (positive) skewed density plot with a long tail. This means that there are quite a few values that sit higher than the mean and that the highest density of values isn't a standardly distributed series. We see that the largest density of purchases is around the 250000 mark.[Figure 11]

4.5 Marital Status

Now examine the marital status of store customers.

```
dataset_maritalStatus = dataset %>%
  select(User_ID, Marital_Status) %>%
  group_by(User_ID) %>%
  distinct()

head(dataset_maritalStatus)
```

User_ID	Marital_Status
1000001	0
1000002	0
1000003	0
1000004	1
1000005	1
1000006	0

Table 17. Marital status
Source: Own elaboration.

Note, we need to quickly change Marital_Status from a numeric variable to a categorical type.

```
dataset_maritalStatus$Marital_Status =
as.character(dataset_maritalStatus$Marital_Status)
typeof(dataset_maritalStatus$Marital_Status)
'character'
```

If we look back at the variable descriptions of the dataset, we don't have a clear guide for marital status. In other cases, it would be best to reach out to the provider of the data to be completely sure of what the values in a column represent but in this case, we will assume that 1 = married and 0 = single.[Table 17]

```
marital_vis = ggplot(data = dataset_maritalStatus) +
  geom_bar(mapping = aes(x = Marital_Status, y =
..count.., fill = Marital_Status)) +
  labs(title = 'Marital Status') +
  scale_fill_brewer(palette = 'Pastel2')
print(marital_vis)
```




Figure 12. Marital status.
Source: Own elaboration.

It looks like most of our shoppers happen to be single or unmarried. Similar to our investigation of age groups, we can look at the makeup of Marital_Status in each City_Category. [Figure 12]

```
dataset_maritalStatus = dataset_maritalStatus %>%
  full_join(customers_stay, by = 'User_ID')
head(dataset_maritalStatus)
```

User_ID	Marital_Status	City_Category	Stay_In_Current_City_Years
1000001	0	A	2
1000002	0	C	4+
1000003	0	A	3
1000004	1	B	2
1000005	1	A	1
1000006	0	A	1

Table 18. Full join (customers_stay, by = 'User_ID'

Source: Own elaboration.

```
maritalStatus_cities = dataset_maritalStatus %>%
  group_by(City_Category, Marital_Status) %>%
  tally()
head(maritalStatus_cities)
```

City_Category	Marital_Status	n
A	0	652
A	1	393
B	0	1004
B	1	703
C	0	1761
C	1	1378

Table.19 City Category

Source: Own elaboration.

```
ggplot(data = maritalStatus_cities, aes(x = City_Category, y = n, fill =
Marital_Status)) +
  geom_bar(stat = "identity", color = 'black') +
  scale_fill_brewer(palette = 2) +
  labs(title = "City + Marital Status",
```



```
y = "Total Count (Shoppers)", x = "City", fill = "Marital Status")
```

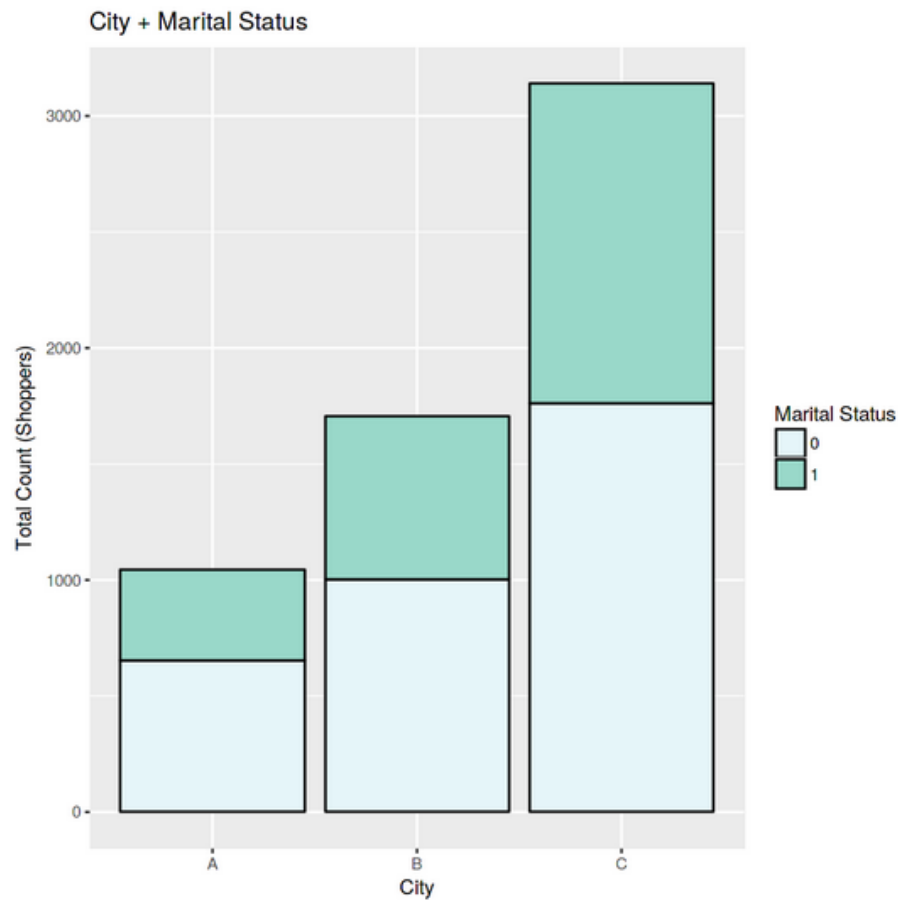


Figure 13. City
Source: Own elaboration.

Here, we can see that out off all Cities, the highest proportion of single shoppers seems to be in City A. Now, lets investigate the Stay_in_Current_City distribution within each City_Category. [Figure 13]

```
Users_Age = dataset %>%
  select(User_ID, Age) %>%
  distinct()
head(Users_Age)
```

User_ID	Age
1000001	0-17
1000002	55+
1000003	26-35
1000004	46-50
1000005	26-35
1000006	51-55

Table 20. ID Age

Source: Own elaboration.

```
dataset_maritalStatus = dataset_maritalStatus %>%
  full_join(Users_Age, by = 'User_ID')
head(dataset_maritalStatus)
```

User_ID	Marital_Status	City_Category	Stay_In_Current_City_Years	Age
1000001	0	A	2	0-17
1000002	0	C	4+	55+
1000003	0	A	3	26-35
1000004	1	B	2	46-50
1000005	1	A	1	26-35
1000006	0	A	1	51-55

Table 21. Data from marital status for User Id
Source: Own elaboration.

```
City_A = dataset_maritalStatus %>%
  filter(City_Category == 'A')
City_B = dataset_maritalStatus %>%
  filter(City_Category == 'B')
City_C = dataset_maritalStatus %>%
  filter(City_Category == 'C')
head(City_A)
head(City_B)
head(City_C)
```

User_ID	Marital_Status	City_Category	Stay_In_Current_City_Years	Age
1000001	0	A	2	0-17
1000003	0	A	3	26-35
1000005	1	A	1	26-35
1000006	0	A	1	51-55
1000015	0	A	1	26-35
1000019	0	A	3	0-17
1000004	1	B	2	46-50
1000007	1	B	1	36-45
1000010	1	B	4+	36-45
1000018	0	B	3	18-25
1000021	0	B	0	18-25
1000023	1	B	3	36-45
1000002	0	C	4+	55+
1000008	1	C	4+	26-35
1000009	0	C	0	26-35
1000011	0	C	1	26-35
1000012	0	C	2	26-35
1000013	1	C	3	46-50

Table 22. Data from City for User Id.
Source: Own elaboration.

```

City_A_stay_vis = ggplot(data = City_A, aes(x = Age, y = ..count.., fill = Age)) +
  geom_bar(stat = 'count') +
  scale_fill_brewer(palette = 8) +
  theme(legend.position="none", axis.text =
element_text(size = 6)) +
  labs(title = 'City A', y = 'Count', x = 'Age', fill =
'Age')
City_B_stay_vis = ggplot(data = City_B, aes(x = Age, y = ..count.., fill = Age)) +
  geom_bar(stat = 'count') +
  scale_fill_brewer(palette = 9) +
  theme(legend.position="none", axis.text =
element_text(size = 6)) +
  labs(title = 'City B', y = 'Count', x = 'Age', fill =
'Age')
City_C_stay_vis = ggplot(data = City_C, aes(x = Age, y = ..count.., fill = Age)) +
  geom_bar(stat = 'count') +
  scale_fill_brewer(palette = 11) +

```

```

element_text(size = 6)) +
  theme(legend.position="none", axis.text =
  labs(title = 'City C', y = 'Count', x = 'Age', fill =
  'Age')
grid.arrange(City_A_stay_vis, City_B_stay_vis, City_C_stay_vis, ncol = 3)

```

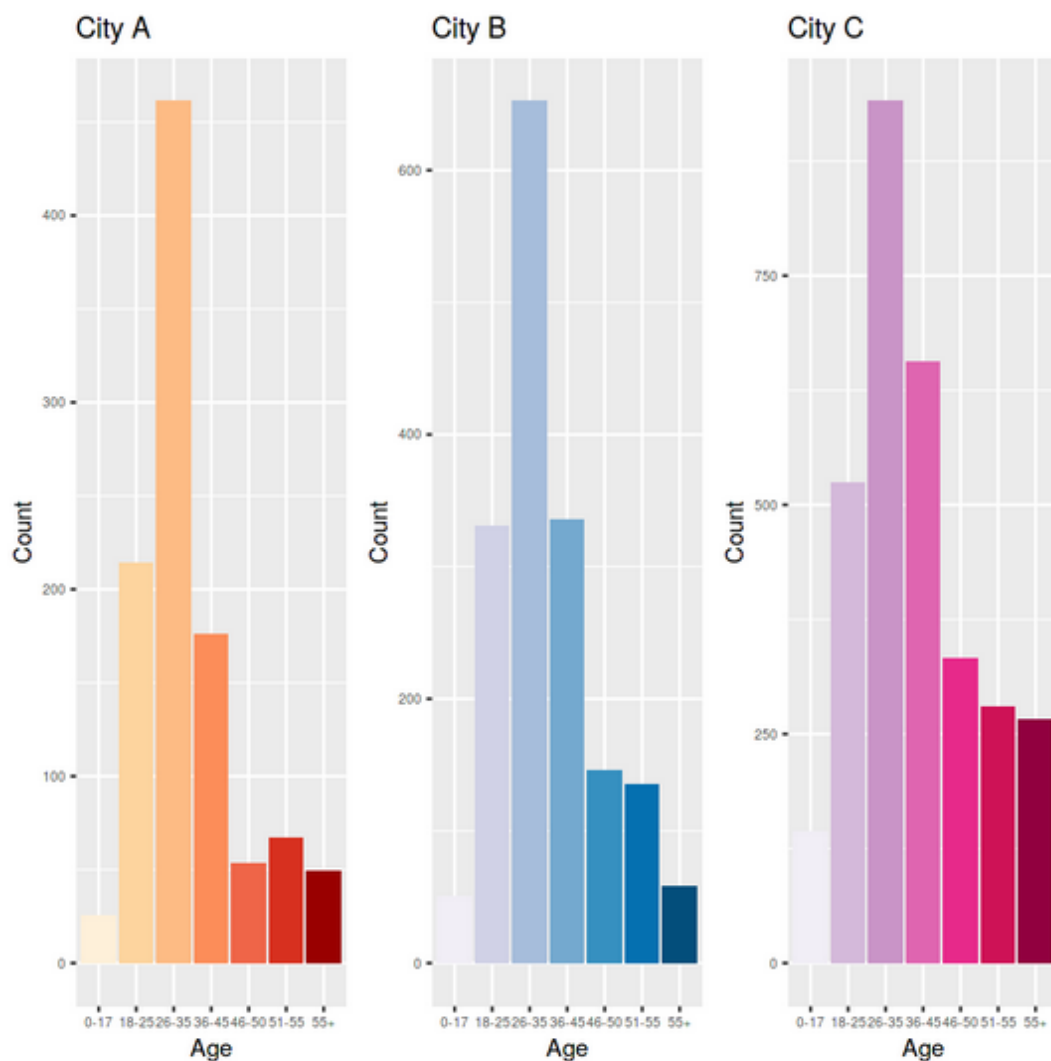


Figure 12. Count for Age
Source: Own elaboration.

It looks as though City A has less shoppers living there over the age of 45 compared to the other cities. This could be a factor in the resulting levels of Marital_Status within each individual city.[Figure 12]

4.5 Top Shoppers

Now we will investigate who our top shoppers were on Black Friday.

```
top_shoppers = dataset %>%  
  count(User_ID, sort = TRUE)  
  
head(top_shoppers)
```

User_ID	n
1001680	1025
1004277	978
1001941	898
1001181	861
1000889	822
1003618	766

Table 23. Data User Id

Source: Own elaboration.

Looks like User_ID 1001680 shows up the most on our master ledger of shopper data. Since each individual row represents a different transaction/product, it looks like this user made over 1000 total transactions! We can join together this top shoppers dataset with our total customer purchases dataset to see them combined.[Table 23]

```
top_shoppers = top_shoppers %>%  
  select(User_ID, n) %>%  
  left_join(customers_total_purchase_amount, Purchase_Amount, by =  
    'User_ID')  
  
head(top_shoppers)
```

User_ID	n	Purchase_Amount
1001680	1025	8699232
1004277	978	10536783
1001941	898	6817493
1001181	861	6387899
1000889	822	5499812
1003618	766	5961987

Table 24. Data User Id Purchase_Amount

Source: Own elaboration.

Now that we have joined the two tables together, we can see that although User_ID 1001680 has the highest number of total purchases, User_ID 1004277 has the highest Purchase_Amount as identified in our earlier charts as well. From here, we can also compute the average Purchase_Amount for each user.[Table 24]

```
top_shoppers = mutate(top_shoppers,
                        Average_Purchase_Amount = Purchase_Amount/n)
head(top_shoppers)
```

User_ID	n	Purchase_Amount	Average_Purchase_Amount
1001680	1025	8699232	8487.056
1004277	978	10536783	10773.807
1001941	898	6817493	7591.863
1001181	861	6387899	7419.163
1000889	822	5499812	6690.769
1003618	766	5961987	7783.273

Table 25. Data User Id Purchase Amount

Source: Own elaboration.

Now, we can sort according to Average_Purchase_Amount to see which customers, on average, are spending the most.[Table 25]

```
top_shoppers_averagePurchase = top_shoppers %>%
                                arrange(desc(Average_Purchase_Amount))
```

```
head(top_shoppers_averagePurchase)
```

User_ID	n	Purchase_Amount	Average_Purchase_Amount
1005069	16	308454	19278.38
1003902	93	1746284	18777.25
1005999	18	330227	18345.94
1001349	23	417743	18162.74
1000101	65	1138239	17511.37
1003461	20	350174	17508.70

Table 26. Data User Id Purchase_Amount Average_Purchase_Amount

Source: Own elaboration.

Looks like User_ID 1005069 has the highest Average_Purchase_Amount and a total Purchase_Amount of 308454. User_ID 1003902 is right behind User_ID 1005069 in Average_Purchase_Amount, but has a much higher total Purchase_Amount of 1746284.

4.6 Occupation

We will analyze is the occupation of customers in our dataset.

```
customers_Occupation = dataset %>%
  select(User_ID, Occupation) %>%
  group_by(User_ID) %>%
  distinct() %>%
  left_join(customers_total_purchase_amount, Occupation, by
= 'User_ID')
head(customers_Occupation)
```

User_ID	Occupation	Purchase_Amount
1000001	10	333481
1000002	16	810353
1000003	15	341635
1000004	7	205987
1000005	20	821001
1000006	9	379450

Table 27. Data Occupation and Purchase Amount

Source: Own elaboration.

Now that we have our dataset necessary, we can group together the total Purchase_Amount for each Occupation identifier. We will then convert Occupation to a character data type [Table 27].

```
totalPurchases_Occupation = customers_Occupation %>%
  group_by(Occupation) %>%
  summarise(Purchase_Amount = sum(Purchase_Amount)) %>%
  arrange(desc(Purchase_Amount))

totalPurchases_Occupation$Occupation =
as.character(totalPurchases_Occupation$Occupation)
typeof(totalPurchases_Occupation$Occupation)

head(totalPurchases_Occupation)
'character'
```

Occupation	Purchase_Amount
4	657530393
0	625814811
7	549282744
1	414552829
17	387240355
12	300672105

Table 28. Data Occupation and Purchase Amount
Source: Own elaboration.

Now, lets plot each occupation and their total [Table 28]Purchase_Amount:

```
occupation = ggplot(data = totalPurchases_Occupation) +
  geom_bar(mapping = aes(x = reorder(Occupation, -Purchase_Amount),
y = Purchase_Amount, fill = Occupation), stat = 'identity') +
  scale_x_discrete(name="Occupation", breaks = seq(0,20, by = 1),
expand = c(0,0)) +
  scale_y_continuous(name="Purchase Amount ($)", expand = c(0,0),
limits = c(0, 750000000)) +
  labs(title = 'Total Purchase Amount by Occupation') +
  theme(legend.position="none")
print(occupation)
```





Figure 13. Total purchase amount by occupation
Source: Own elaboration.

Looks like customers labeled as Occupation 4 spent the most at our store on Black Friday, with customers of Occupation 0 + 7 closely behind. Here, if a key was given, we could use that information to classify our shoppers accordingly.[Figure 13]

4.7 Apriori (Association Rule Learning)

Apriori is an algorithm for frequent item set mining and association rule learning over transactional databases. It proceeds by identifying the frequent individual items in the database and extending them to larger and larger item sets as long as those item sets appear sufficiently often in the database. The frequent item sets determined by Apriori can be used to determine association rules which highlight general trends in the database: this has applications in domains such as market basket analysis.¹¹

Now lets use a machine learning algorithm called Apriori to make some association rules regarding customer purchases. We will be using the arules package.

Before we begin, lets elaborate on the idea of Association Rule Learning. In its simplest form, Association Rule Learning attempts to predict customer transactions. In other words, the algorithm

¹¹ https://en.wikipedia.org/wiki/Apriori_algorithm

solves the problem, "People who bought ----- also bought ----- ." This can prove to be extremely useful for retailers who aim to optimize product placement in stores and promotional campaigns.¹²

In the case of our store on Black Friday, implementing an effective product placement strategy can prove to optimize sales of products normally bought together. For example, let's say that our store was to have a sale on TVs. It would be smart to place HDMI Cables alongside these TVs because those items are usually purchased together. On the other hand, it may also prove to be smart to place them far apart so that customers need to walk throughout the entire store while searching for their desired item, where another product may catch their eye along the way.

The Apriori algorithm specifically aims to maximize the likelihood someone performs/purchases something given knowledge about their prior actions.

To begin, let's import the libraries we will be using for this section if not done so already.

```
library(arules)
library(arulesViz)
library(tidyverse)
Loading required package: grid
```

The arules package was developed specifically to deal with Association Rule and Frequent Itemset mining. In order to begin our analysis, we must retrieve the necessary data from the original dataset and then apply the correct formatting.

```
# Data Preprocessing
# Getting the dataset into the correct format
customers_products = dataset %>%
  select(User_ID, Product_ID) %>% # Selecting the columns we
will need
  group_by(User_ID) %>% # Grouping by "User_ID"
  arrange(User_ID) %>% # Arranging by "User_ID"
  mutate(id = row_number()) %>% # Defining a key column for
each "Product_ID" and its corresponding "User_ID" (Must do this for spread() to work
properly)
  spread(User_ID, Product_ID) %>% # Converting our dataset
from tall to wide format, and grouping "Product_IDs" to their corresponding "User_ID"
  t() # Transposing the dataset
from columns of "User_ID" to rows of "User_ID"

# Now we can remove the Id row we created earlier for spread() to work correctly.
customers_products = customers_products[-1,]
```

Now, in order for the Apriori algorithm to work correctly, we need to convert the customers_products table into a sparse matrix. Unfortunately, Apriori doesn't take strings or text as input, but rather 1 + 0. (Binary Format) This means that we must allocate a column for each individual product and then if a

¹² https://en.wikipedia.org/wiki/Apriori_algorithm

User_ID contains that product, it will be marked as a 1. On the other hand, if the User_ID does not contain that Product_ID, it will be marked with a 0.

In order to do so, we need to use the arules library as described above and import the table as a .csv file. From there, we can use the arules function, "read.transactions()" to get our sparse matrix.

```
write.csv(customers_products, file = 'customers_products.csv')

customersProducts = read.transactions('customers_products.csv', sep = ',',
rm.duplicates = TRUE) # remove duplicates with rm.duplicates
distribution of transactions with duplicates:
```

items	46	126	163	202	258	272	285	307	310	316	319	327	330	334	340	344
	1	1	1	1	1	1	1	1	1	1	1	2	1	1	1	1
345	348	354	357	373	393	402	408	419	437	441	449	450	452	454	456	
	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	
459	465	466	467	475	476	477	481	487	491	495	498	507	523	524	526	
	1	1	2	2	1	1	1	2	2	3	2	1	1	2	1	
527	528	530	531	532	533	535	537	538	539	540	545	546	548	549	553	
	2	1	1	2	1	1	1	2	3	1	1	1	3	1	1	
554	555	556	558	563	566	567	570	572	574	575	577	578	580	583	584	
	2	1	1	2	1	3	1	3	2	4	1	2	3	2	1	
586	588	589	590	591	592	593	594	595	597	598	601	602	604	607	608	
	2	3	5	1	1	1	1	3	2	1	1	1	1	2	1	
610	612	613	614	615	616	617	618	619	620	623	625	632	633	634	635	
	1	2	1	2	1	1	2	1	3	2	1	1	6	1	5	
638	640	641	642	643	644	645	646	647	648	653	654	657	658	659	661	
	2	2	1	2	1	2	3	1	4	1	1	3	2	1	1	
662	663	664	665	666	667	668	669	670	671	672	674	676	677	678	679	
	1	3	1	1	2	1	2	4	2	1	3	1	1	2	3	
681	682	683	685	686	687	688	689	690	691	692	694	695	697	698	699	
	2	2	4	4	5	2	2	1	1	1	2	1	1	1	1	
700	702	703	704	705	706	707	708	709	710	712	713	714	715	716	717	
	2	1	3	1	1	2	2	3	2	2	4	5	2	2	1	
718	719	720	721	722	723	724	725	726	727	728	729	730	732	733	734	
	2	3	3	5	2	1	1	2	5	2	1	3	3	2	1	
735	736	737	738	739	740	741	742	743	744	745	747	748	749	750	751	
	6	2	4	6	3	4	1	6	7	4	6	1	1	5	5	
752	753	754	755	756	757	758	759	760	761	763	764	765	766	767	768	
	2	3	4	6	6	2	6	2	1	5	3	4	2	3	2	
769	770	771	772	773	774	775	776	777	778	779	780	781	782	783	784	
	2	2	3	1	3	4	2	2	3	3	2	4	7	3	4	
785	786	787	788	789	790	791	792	793	794	795	796	797	798	799	800	
	5	3	3	6	5	5	2	3	5	3	8	5	5	9	3	
801	802	803	804	805	806	807	808	809	810	811	812	813	814	815	816	
	4	7	4	3	4	5	7	5	4	5	3	2	6	6	3	
817	818	819	820	821	822	823	824	825	826	827	828	829	830	831	832	
	5	10	6	6	4	7	7	2	5	4	7	5	5	4	5	
833	834	835	836	837	838	839	840	841	842	843	844	845	846	847	848	
	3	5	4	11	5	5	4	9	7	6	4	7	9	11	4	
849	850	851	852	853	854	855	856	857	858	859	860	861	862	863	864	
	10	6	10	7	12	16	11	8	7	4	12	9	11	11	9	
865	866	867	868	869	870	871	872	873	874	875	876	877	878	879	880	
	11	10	7	6	5	12	6	7	8	11	9	9	8	7	5	
881	882	883	884	885	886	887	888	889	890	891	892	893	894	895	896	
	15	13	12	8	4	6	12	15	13	10	11	13	6	21	7	
897	898	899	900	901	902	903	904	905	906	907	908	909	910	911	912	
	9	7	11	18	5	14	10	9	19	15	10	17	18	23	8	
913	914	915	916	917	918	919	920	921	922	923	924	925	926	927	928	
	15	12	18	21	17	12	11	13	13	12	20	20	16	13	15	
929	930	931	932	933	934	935	936	937	938	939	940	941	942	943	944	
	27	22	20	28	18	14	20	20	20	14	22	30	23	23	21	
945	946	947	948	949	950	951	952	953	954	955	956	957	958	959	960	

25	19	30	31	30	24	27	25	40	30	31	16	29	30	32	48
961	962	963	964	965	966	967	968	969	970	971	972	973	974	975	976
27	27	24	30	26	35	43	30	51	49	40	41	36	32	36	38
977	978	979	980	981	982	983	984	985	986	987	988	989	990	991	992
43	41	42	37	49	44	51	57	55	40	53	56	63	39	58	50
993	994	995	996	997	998	999	1000	1001	1002	1003	1004	1005	1006	1007	1008
58	77	74	72	72	84	74	66	77	85	93	79	94	118	122	104
1009	1010	1011	1012	1013	1014	1015	1016	1017	1018	1019					
121	113	120	78	77	55	37	20	7	5	1					

Table 29. Sparse matrix customers products.
Source: Own elaboration.

Before we implement the Apriori algorithm to our problem, lets take a look at our newly created sparse matrix.[Table 29]

In numerical analysis and scientific computing, a sparse matrix or sparse array is a matrix in which most of the elements are zero. By contrast, if most of the elements are nonzero, then the matrix is considered dense. The number of zero-valued elements divided by the total number of elements (e.g., $m \times n$ for an $m \times n$ matrix) is called the sparsity of the matrix (which is equal to 1 minus the density of the matrix). Using those definitions, a matrix will be sparse when its sparsity is greater than 0.5.¹³

```
summary(customersProducts)
transactions as itemMatrix in sparse format with
  5892 rows (elements/itemsets/transactions) and
  10539 columns (items) and a density of 0.008768598

most frequent items:
P00265242 P00110742 P00025442 P00112142 P00057642 (Other)
  1858      1591      1586      1539      1430      536489

element (itemset/transaction) length distribution:
sizes
  6    7    8    9   10   11   12   13   14   15   16   17   18   19   20   21
  1    5    7   20   37   55   77   78  120  113  121  104  122  118   94   79
 22   23   24   25   26   27   28   29   30   31   32   33   34   35   36   37
 93   85   77   66   74   84   72   72   74   77   58   50   58   39   63   56
 38   39   40   41   42   43   44   45   46   47   48   49   50   51   52   53
 53   40   55   57   51   44   49   37   42   41   43   38   36   32   36   41
 54   55   56   57   58   59   60   61   62   63   64   65   66   67   68   69
 40   49   51   30   43   35   26   30   24   27   27   48   32   30   29   16
 70   71   72   73   74   75   76   77   78   79   80   81   82   83   84   85
 31   30   40   25   27   24   30   31   30   19   25   20   21   23   23   30
 86   87   88   89   90   91   92   93   94   95   96   97   98   99  100  101
 22   14   20   20   20   14   18   28   20   22   27   17   15   13   16   20
102  103  104  105  106  107  108  109  110  111  112  113  114  115  116  117
 20   12   13   13   11   12   17   21   18   12   15   19   8    23   18   17
118  119  120  121  122  123  124  125  126  127  128  129  130  131  132  133
 10   15   19   9    10   14   5    18   11   7    9   14   7    21   6    13
134  135  136  137  138  139  140  141  142  143  144  145  146  147  148  149
 11   10   13   15   12   6    4    8   12   13   15   4    5    7    8    9
150  151  152  153  154  155  156  157  158  159  160  161  162  163  164  165
 9    11   8    7    6    12   5    6    7   10   11   6    9   11   11   9
166  167  168  169  170  171  172  173  174  175  176  177  178  179  180  181
 12   4    7    8   11   16   12   7   10   6   10   6    4   11   9    7
```

¹³ https://en.wikipedia.org/wiki/Sparse_matrix

182	183	184	185	186	187	188	189	190	191	192	193	194	195	196	197
4	6	7	9	4	5	5	11	4	5	3	4	5	4	5	5
198	199	200	201	202	203	204	205	206	207	208	209	210	211	212	213
7	4	5	2	7	7	4	6	6	10	5	11	3	6	6	2
214	215	216	217	218	219	220	221	222	223	224	225	226	227	228	229
3	5	4	5	7	5	4	3	4	7	4	4	3	9	5	5
230	231	232	233	234	235	236	237	238	239	240	241	242	243	244	245
8	3	5	3	2	5	5	6	3	3	5	5	4	3	7	4
246	247	248	249	250	251	252	253	254	255	256	257	258	259	260	261
2	3	3	2	2	4	3	1	3	2	2	5	2	3	2	4
262	264	265	266	267	268	269	270	271	272	273	274	275	276	277	278
3	5	1	2	6	2	6	6	4	3	2	3	5	5	1	1
280	281	282	283	284	285	286	287	288	289	290	291	292	293	295	296
6	4	7	6	1	4	3	6	4	2	6	5	1	2	3	3
297	298	299	300	301	302	303	304	305	306	307	308	309	310	311	312
1	2	5	2	1	1	2	5	3	3	2	1	1	2	2	5
313	315	316	317	318	319	320	321	322	323	325	326	327	328	330	331
4	2	2	3	2	2	1	1	3	1	2	1	1	1	1	2
333	334	335	336	337	338	339	340	342	343	344	346	347	348	349	351
1	1	1	1	2	2	5	4	4	2	2	3	3	2	1	1
353	354	355	356	357	358	359	360	361	362	363	364	366	367	368	371
3	1	2	4	2	1	2	1	1	3	1	2	1	1	2	3
372	377	378	379	380	381	382	383	384	385	387	390	391	392	393	400
1	1	4	1	3	2	1	2	1	2	2	2	5	1	6	1
402	405	406	407	408	409	410	411	412	413	415	417	418	421	423	424
1	2	3	1	2	1	1	2	1	2	1	1	2	1	1	1
427	428	430	431	432	433	434	435	436	437	439	441	442	445	447	448
1	2	3	1	1	1	1	1	5	3	2	1	1	2	3	2
450	451	453	455	458	459	462	467	469	470	471	472	476	477	479	480
1	4	2	3	1	3	1	2	1	1	2	1	1	3	1	1
485	486	487	488	490	492	493	494	495	497	498	499	501	502	518	527
1	3	2	1	1	1	1	2	1	1	2	1	2	1	1	2
530	534	538	544	548	549	550	558	559	560	566	569	571	573	575	576
3	2	2	1	1	1	1	2	2	1	1	1	1	1	1	1
584	588	606	617	623	632	652	668	671	677	680	681	685	691	695	698
1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	2
706	709	715	718	740	753	767	823	862	899	979	1025	1026			
1	1	1	1	1	1	1	1	1	1	1	1	1			

Min. 1st Qu. Median Mean 3rd Qu. Max.
6.00 26.00 54.00 92.41 115.00 1026.00

Table 30. Transactions as item Matrix.
Source: Own elaboration.

Here, we can see that there are 5892 rows (elements/itemsets/transactions) and 10539 columns (items) in our sparse matrix. With this summary function, we get a density of 0.008768598 in our matrix. The density tells us that we have 0.9% non-zero values (1) in our sparse matrix and 99.1% zero (0) values.

Also, as we discovered in our Exploratory Data Analysis, the summary() function also gives us the most frequent items that customers purchased and just to be sure, we can cross reference what we discovered earlier in the analysis. Lets list out what our sparse matrix gave us.[Table 30]

- P00265242 = 1858
- P00110742 = 1591
- P00025442 = 1586

- P00112142 = 1539
- P00057642 = 1430
- (Other) = 536489

Now we can compare it to what we discovered earlier.

"Looks like our top 5 best sellers are (by product ID)"

- P00265242 = 1858
- P00110742 = 1591
- P00025442 = 1586
- P00112142 = 1539
- P00057642 = 1430

Looks like our sparse matrix is accurate to what we discovered earlier. It is important to ensure that all data is being transferred correctly in every step of the analysis. This ensures repeatability and easy debugging should an error occur.

Continue to examine our sparse matrix.

```
summary(customersProducts)
transactions as itemMatrix in sparse format with
  5892 rows (elements/itemsets/transactions) and
  10539 columns (items) and a density of 0.008768598
```

```
most frequent items:
P00265242 P00110742 P00025442 P00112142 P00057642 (Other)
  1858      1591      1586      1539      1430      536489
```

```
element (itemset/transaction) length distribution:
sizes
```

6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21
1	5	7	20	37	55	77	78	120	113	121	104	122	118	94	79
22	23	24	25	26	27	28	29	30	31	32	33	34	35	36	37
93	85	77	66	74	84	72	72	74	77	58	50	58	39	63	56
38	39	40	41	42	43	44	45	46	47	48	49	50	51	52	53
53	40	55	57	51	44	49	37	42	41	43	38	36	32	36	41
54	55	56	57	58	59	60	61	62	63	64	65	66	67	68	69
40	49	51	30	43	35	26	30	24	27	27	48	32	30	29	16
70	71	72	73	74	75	76	77	78	79	80	81	82	83	84	85
31	30	40	25	27	24	30	31	30	19	25	20	21	23	23	30
86	87	88	89	90	91	92	93	94	95	96	97	98	99	100	101
22	14	20	20	20	14	18	28	20	22	27	17	15	13	16	20
102	103	104	105	106	107	108	109	110	111	112	113	114	115	116	117
20	12	13	13	11	12	17	21	18	12	15	19	8	23	18	17
118	119	120	121	122	123	124	125	126	127	128	129	130	131	132	133
10	15	19	9	10	14	5	18	11	7	9	14	7	21	6	13
134	135	136	137	138	139	140	141	142	143	144	145	146	147	148	149
11	10	13	15	12	6	4	8	12	13	15	4	5	7	8	9
150	151	152	153	154	155	156	157	158	159	160	161	162	163	164	165
9	11	8	7	6	12	5	6	7	10	11	6	9	11	11	9
166	167	168	169	170	171	172	173	174	175	176	177	178	179	180	181
12	4	7	8	11	16	12	7	10	6	10	6	4	11	9	7

182	183	184	185	186	187	188	189	190	191	192	193	194	195	196	197
4	6	7	9	4	5	5	11	4	5	3	4	5	4	5	5
198	199	200	201	202	203	204	205	206	207	208	209	210	211	212	213
7	4	5	2	7	7	4	6	6	10	5	11	3	6	6	2
214	215	216	217	218	219	220	221	222	223	224	225	226	227	228	229
3	5	4	5	7	5	4	3	4	7	4	4	3	9	5	5
230	231	232	233	234	235	236	237	238	239	240	241	242	243	244	245
8	3	5	3	2	5	5	6	3	3	5	5	4	3	7	4
246	247	248	249	250	251	252	253	254	255	256	257	258	259	260	261
2	3	3	2	2	4	3	1	3	2	2	5	2	3	2	4
262	264	265	266	267	268	269	270	271	272	273	274	275	276	277	278
3	5	1	2	6	2	6	6	4	3	2	3	5	5	1	1
280	281	282	283	284	285	286	287	288	289	290	291	292	293	295	296
6	4	7	6	1	4	3	6	4	2	6	5	1	2	3	3
297	298	299	300	301	302	303	304	305	306	307	308	309	310	311	312
1	2	5	2	1	1	2	5	3	3	2	1	1	2	2	5
313	315	316	317	318	319	320	321	322	323	325	326	327	328	330	331
4	2	2	3	2	2	1	1	3	1	2	1	1	1	1	2
333	334	335	336	337	338	339	340	342	343	344	346	347	348	349	351
1	1	1	1	2	2	5	4	4	2	2	3	3	2	1	1
353	354	355	356	357	358	359	360	361	362	363	364	366	367	368	371
3	1	2	4	2	1	2	1	1	3	1	2	1	1	2	3
372	377	378	379	380	381	382	383	384	385	387	390	391	392	393	400
1	1	4	1	3	2	1	2	1	2	2	2	5	1	6	1
402	405	406	407	408	409	410	411	412	413	415	417	418	421	423	424
1	2	3	1	2	1	1	2	1	2	1	1	2	1	1	1
427	428	430	431	432	433	434	435	436	437	439	441	442	445	447	448
1	2	3	1	1	1	1	1	5	3	2	1	1	2	3	2
450	451	453	455	458	459	462	467	469	470	471	472	476	477	479	480
1	4	2	3	1	3	1	2	1	1	2	1	1	3	1	1
485	486	487	488	490	492	493	494	495	497	498	499	501	502	518	527
1	3	2	1	1	1	1	2	1	1	2	1	2	1	1	2
530	534	538	544	548	549	550	558	559	560	566	569	571	573	575	576
3	2	2	1	1	1	1	2	2	1	1	1	1	1	1	1
584	588	606	617	623	632	652	668	671	677	680	681	685	691	695	698
1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	2
706	709	715	718	740	753	767	823	862	899	979	1025	1026			
1	1	1	1	1	1	1	1	1	1	1	1	1			
Min.		1st Qu.	Median		Mean		3rd Qu.	Max.							
6.00		26.00	54.00		92.41		115.00	1026.00							

Table 31. Transactions as item Matrix.
Source: Own elaboration.

The "element (itemset/transaction) length distribution" gives us a distribution of the number of items in a customers (User) basket and underneath it we can see more information including the quartile and mean information. In this case, we see a mean of 92.41, which means that on average, each customer purchased 92.41 items. In this case, since we are aware of a few customers who purchased over ~1000 items, it may be useful to use the median value of 54.00 items instead since the mean can be heavily affected by outlier values.

To get a clearer picture of the items, lets create an item frequency plot which is included in the arules package.[Table 31]

```
itemFrequencyPlot(customersProducts, topN = 25)    # topN is limiting to the top 50
products
```




Figure 14. Frequency Plot customers Products
Source: Own elaboration.

Now we begin training the association rule model.

Our first step will be to set our parameters. The first parameters we will set are the support and confidence. The support value is derived from the frequency of a specific item within the dataset. When we set our support value, we are setting a minimum number of transactions necessary for our rules to take effect.

- Support: Our support value will be the minimum number of transactions necessary divided by the total number of transactions.

- As described by summary(customersProducts), we have a total number of unique customer transactions of 5892.
- From our dataset, let's assume that we want to choose a product which was purchased by at least 50 different customers.
- With these two values established, we can compute the support value with simple division. $(50/5892) = .008486083$

The second parameter we will take into consideration will be the confidence. The confidence value determines how often a rule is to be found true. In other words, the minimum strength of any rule is a limit we place when setting our minimum confidence value.

The default confidence value in the apriori() function is 0.80 or 80%, so we can begin with that number and then adjust the parameters to applicable results.

- Confidence: We can determine our confidence value by first starting with the default value and adjusting accordingly.
 - With more domain knowledge, and with Product_IDs referencing items with recognizable names, the Confidence value can be easily changed to see different, and more relevant, results.
 - In our case, we will start with a value and then lower the confidence to see different rules.

```
rules = apriori(data = customersProducts,
                parameter = list(support = 0.008, confidence = 0.80, maxtime = 0)) #
maxtime = 0 will allow our algorithm to run until completion with no time limit
Apriori
```

Parameter specification:

```
confidence minval smax arem aval originalSupport maxtime support minlen
0.8 0.1 1 none FALSE TRUE 0 0.008 1
maxlen target ext
10 rules FALSE
```

Algorithmic control:

```
filter tree heap memopt load sort verbose
0.1 TRUE TRUE FALSE TRUE 2 TRUE
```

Absolute minimum support count: 47

```
set item appearances ...[0 item(s)] done [0.00s].
set transactions ...[10539 item(s), 5892 transaction(s)] done [0.10s].
sorting and recoding items ... [2099 item(s)] done [0.02s].
creating transaction tree ... done [0.00s].
checking subsets of size 1 2 3 4 5 6 done [19.52s].
writing ... [7 rule(s)] done [0.47s].
creating S4 object ... done [0.24s].
```

It looks like apriori has created 7 rules in accordance to our specified parameters.

```
"writing ... [7 rule(s)] done [0.48s]."
```

Now we can examine our results to get a better idea of how our algorithm worked.

```
inspect(sort(rules, by = 'lift'))
lhs      rhs      support confidence    lift count
[1] {P00032042,
P00057642,
P00102642,
P00145042} => {P00270942} 0.008655804  0.8793103 4.540663    51
[2] {P00025442,
P00031042,
P00034742,
P00255842} => {P00145042} 0.008486083  0.8064516 3.433246    50
[3] {P00003242,
P00130742,
P00237542} => {P00145042} 0.008316361  0.8032787 3.419738    49
[4] {P00006942,
P00251242,
P00277642} => {P00145042} 0.009674134  0.8028169 3.417773    57
[5] {P00034042,
P00112442,
P00112542} => {P00110742} 0.008146640  0.8135593 3.012880    48
[6] {P00127642,
P00165442,
P00277442} => {P00110742} 0.008316361  0.8032787 2.974807    49
[7] {P00051442,
P00112142,
P00112542,
P00270942} => {P00110742} 0.008146640  0.8000000 2.962665    48
```

Table 32. Transactions apriori for data customersProducts
Source: Own elaboration.

We present the association rules created by our apriori algorithm. Let's take a look at rule number 1.

We see a few values listed and we will go through them individually.

- The first value lhs, corresponds to a grouping of items which the algorithm has pulled from the dataset.
- The second value, rhs, corresponds to the value predicted by apriori to be purchased with items in the "lhs" category.
- The third value, support is the number of transactions including that specific set of items divided by the total number of transactions. (As described earlier when we chose the parameters for Apriori.)
- The fourth value, confidence is the % chance in which a rule will be upheld.
- The fifth value, lift gives us the independance/dependence of a rule. It takes the confidence value and its relationship to the entire dataset into account.
- The sixth and final value, count is the number of times a rule occurred during the implementation of Apriori on our data.

Now, lets visualize these rules using the arulesViz package.

```
plot(rules, method = 'graph')
```




Figure 15. Graph present 7 rules.
Source: Own elaboration.

In a visualization of our association rules. Arrows pointing from items to rule vertices indicate LHS (Grouped) items and arrows from rules to items indicates the RHS (Rule Item).

The size of the bubbles indicate the support with larger bubbles representing a higher support value. Fill color represents the lift values, with darker colors representing higher lifts.

Lets now try modifying some of the parameters for the Apriori algorithm and see the results. This process would prove to be more intuitive if given a key for each corresponding Product_ID, so will only implement the algorithm once more.

This time, we will decrease our confidence value to 75% and keep our support value the same (0.008).

```
ules = apriori(data = customersProducts,
               parameter = list(support = 0.008, confidence = 0.75, maxtime = 0))
Apriori
```

```

Parameter specification:
confidence minval smax arem aval originalSupport maxtime support minlen
      0.75      0.1      1 none FALSE              TRUE          0      0.008      1
maxlen target   ext
      10      rules FALSE

```

```

Algorithmic control:
filter tree heap memopt load sort verbose
      0.1 TRUE TRUE  FALSE TRUE      2      TRUE

```

Absolute minimum support count: 47

```

set item appearances ...[0 item(s)] done [0.00s].
set transactions ...[10539 item(s), 5892 transaction(s)] done [0.10s].
sorting and recoding items ... [2099 item(s)] done [0.02s].
creating transaction tree ... done [0.00s].
checking subsets of size 1 2 3 4 5 6 done [19.37s].
writing ... [171 rule(s)] done [0.46s].
creating S4 object ... done [0.25s].

```

Now that we have decreased the minimum confidence value to 75%, we have a total of 171 rules.

```
writing ... [171 rule(s)] done [0.50s].
```

This is much higher number of rules compared to our previous rule list which only contained 7. This should now give us more interesting rules to examine.

```

inspect(head(sort(rules, by = 'lift')) # limiting to the top 6 rules
lhs      rhs      support confidence      lift count
[1] {P00221142,
P00249642} => {P00103042} 0.008146640  0.7619048 8.030667      48
[2] {P00002142,
P00103042,
P00147942} => {P00221442} 0.008146640  0.7500000 6.045144      48
[3] {P00032042,
P00057642,
P00102642,
P00145042} => {P00270942} 0.008655804  0.8793103 4.540663      51
[4] {P00062842,
P00127242,
P00243942} => {P00044442} 0.008486083  0.7575758 4.061544      50
[5] {P00030842,
P00057942,
P00355142} => {P00114942} 0.008486083  0.7936508 4.024260      50
[6] {P00030842,
P00147742,
P00303342} => {P00044442} 0.008146640  0.7500000 4.020928      48

```

Table 33. Limiting to the top 6 rules●

Source: Own elaboration.

We can now see that we now have a new set of rules and the rule with the highest lift value has also changed.

Rule number 1 shows that Customers who bought items P00221142 and P00249642 will also purchase item P00103042 ~76% of the time, given a support of 0.008.

```
plot(rules, method = 'graph', max = 25)
```

Warning message:

"plot: Too many rules supplied. Only plotting the best 25 rules using 'support' (change control parameter max if needed)"




Figure 16. Graph present the 25 rules
Source: Own elaboration.

Now that we have more than 7 rules, this visualization becomes a lot more difficult to interpret.

Instead, we can create a matrix and have a similar plot and clearer interpretation.

```
plot(rules, method = 'grouped', max = 25)
```

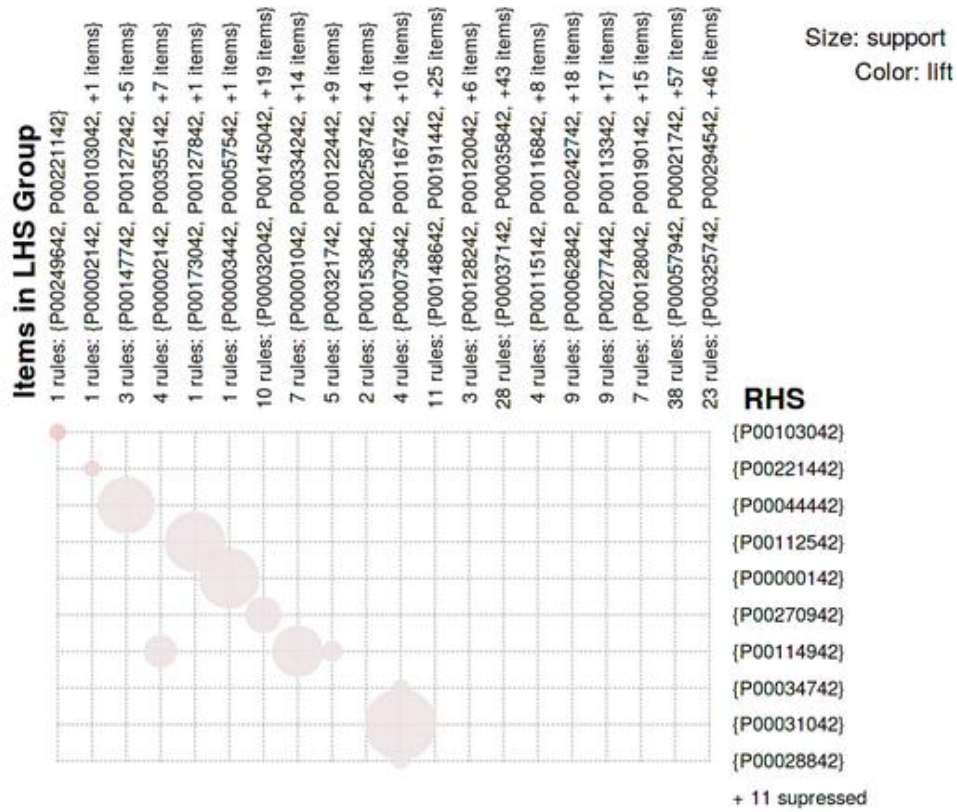


Figure 17. Graph LHS.
 Source: Own elaboration.

In this visualization, we can see that we have our LHS on top and on the right hand side, the corresponding RHS. The size of the bubbles represents the support value of the rule and the fill / color represents the lift.

5. Conclusions

1. After the analysis of this data, we can confirm the male customers have a higher average spending than the female.
2. We can see the gender of the customer are we rejection because the most of customers are men.
3. The category of age of customers we can see is 26-35. From this data and analyses now we can confirm the hypotheses.

References

1. Chang, E., Burns, L. D., & Francis, S. K. (2004) (Abstract)
2. Oxford Dictionaries- Black Friday
3. [https://en.wikipedia.org/wiki/Black_Friday_\(shopping\)](https://en.wikipedia.org/wiki/Black_Friday_(shopping))
4. <https://www.kaggle.com/mehdidag/black-friday>
5. <https://en.wikipedia.org/wiki/Ggplot2>
6. <https://www.history.com/news/whats-the-real-history-of-black-friday>
7. <https://en.oxforddictionaries.com/explore/why-is-day-after-thanksgiving-black-friday/>
8. <https://www.cnn.com/2018/11/21/business/black-friday-history/index.html>
9. <https://journals.sagepub.com/doi/abs/10.1177/0887302X0402200404#articleCitationDownloadContainer>
10. https://en.wikipedia.org/wiki/Apriori_algorithm
11. https://en.wikipedia.org/wiki/Association_rule_learning

Packages

1. <https://www.tidyverse.org/>
2. <https://cran.r-project.org/web/packages/scales/scales.pdf>
3. <https://cran.r-project.org/web/packages/arules/arules.pdf>
4. <https://cran.r-project.org/web/packages/plot.ly/r/>
5. <https://www.numpy.org/>
6. <https://pandas.pydata.org/>
7. <https://matplotlib.org/>
8. [ges/arulesViz/vignettes/arulesViz.pdf](https://arulesViz.vignettes/arulesViz.pdf)

Images

https://justcreative.com/wp-content/uploads/2018/10/black-friday-deals.jpg	2
---	---

Tables

Table 1. Code from R program with data analyses.....	8
Table 2. Information about columns data.	9
Table 3. First 10 data of dataframe to have knowledge data itself.	9
Table 4. First 10 data of dataframe to have knowledge data itself.	10
Table 5. First 7 data of dataframe to have knowledge data itself.	11
Table 6. Overview of the entire dataset.	11
Table 7. Overview of the entire dataset.	12
Table 8. Dupliptes.	13
Table 9. F and M in data Dupliptes.	14
Table 11. Total Purchase.	15
Table 11a. Total Purchase.	16
Table 12. Join Table 2 and Table 3 together.	16
Table 13. Join Purchase and Average Together.	16
Table 13. 5 best selling products.	16
Table 13. Age.	23
Table 14. Purchase Amount.	27
Table 15. Purchase Amount.	28
Table 16. Customers total purchase amount.	28
Table 17. Marital status.	30
Table 18. Full join (customers_stay, by = 'User_ID'.	32
Table.19. City Category.	32
Table 20. ID Age.	33
Table 21. Data prom marital status for User Id.	34
Table 22. Data prom City for User Id.	35
Table 23. Data User Id.	37
Table 24. Data User Id Purchase_Amount.	38
Table 25. Data User Id Purchase Amount.	38
Table 26. Data User Id Purchase_Amount Average_Purchase_Amount.	39
Table 27. Data Occupation and Purchase Amount.	39

Table 28. Data Occupation and Purchase Amount.	40
Table 30. Transactions as item Matrix.	45
Table 31. Transactions as item Matrix.	47
Table 32. Transactions apriori for data customersProducts.	46
Table 33. Limiting to the top 6 rules.	50

Figures

Figure 1. Correlation matrix.....	14
Figure 2. Gender plot.....	18
Figure 3. Average spending by Gender plot.	21
Figure 4. Count and Gender plot.....	22
Figure 6. Proportion of customers.....	25
Figure 7. Age of customers.....	26
Figure 8. Age of customers best seller.....	29
Figure 9. Age of customers best seler.....	31
Figure 10. Age of customers best seller.....	32
Figure 11. Purchase Amount to Density.....	33
Figure 12. Marital status.	36
Figure 13. City.....	41
Figure 14. Frequency Plot customers Products.....	48
Figure 15. Graph present 7 rules.....	51
Figure 15. Graph present 7 rules.....	52
Figure 16. Graph present the 25 rules.....	53
Figure 17. Graph LHS.	54